



Networking

Service discovery

- Kubernetes is quite dynamic with placing pods on nodes, horizontal autoscalers etc.
- Service discovery helps us find the services in the cluster
- Should resolve quickly and reliably
-

The Service Object

A Service Object is a way to create a named label selector

Creates a service:

```
kubectl expose deployment nginx
```

Show including the selector:

```
kubectl get services -o wide
```

```
PS >kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
default-http-backend	ClusterIP	10.152.183.94	<none>	80/TCP	45h	app=default-http-backend
kubernetes	ClusterIP	10.152.183.1	<none>	443/TCP	4d20h	<none>
webapp	ClusterIP	10.152.183.44	<none>	8080/TCP	44h	run=webapp

Port-forward to service

```
# Connect to the pod
$NGINX_POD=$(kubectl get pods -l run=webapp \
-o jsonpath='{.items[0].metadata.name}')
```

Service DNS

- Kubernetes provides a DNS service exposed to Pods
- System component running in and managed by K8s

```
kubectl exec -ti busybox -- nslookup webapp
```

```
PS >kubectl exec -ti busybox -- nslookup webapp
Server:      10.152.183.10
Address 1: 10.152.183.10 kube-dns.kube-system.svc.cluster.local

Name:        webapp
Address 1: 10.152.183.44 webapp.default.svc.cluster.local
```

- Kubectl edit service webapp
- Change type to NodePort
- Kubectl describe service webapp

```
PS >kubectl describe service webapp
Name:                        webapp
Namespace:                  default
Labels:                    run=webapp
Annotations:               <none>
Selector:                  run=webapp
Type:                      NodePort
IP:                        10.152.183.44
Port:                      <unset> 8080/TCP
TargetPort:                8080/TCP
NodePort:                  <unset> 30430/TCP
Endpoints:                10.1.1.26:8080
Session Affinity:          None
External Traffic Policy:   Cluster
Events:                   <none>
```

LoadBalancer

- Only available if cloud supports it
 - Builds on NodePorts
 - Creates a loadbalancer in the cloud and direct it at nodes in cluster
-
- Kubectl edit service webapp
 - Change spec.type to LoadBalancer
 - Kubectl get services now gets an external ip (pending)


Endpoints

- Endpoints are created as a pair with services
- Kubectl describe endpoints webapp

```
PS >kubectl describe endpoints webapp
Name:          webapp
Namespace:     default
Labels:        run=webapp
Annotations:    <none>
Subsets:
  Addresses:          10.1.1.26
  NotReadyAddresses:  <none>
  Ports:
    Name      Port  Protocol
    ----      -
    <unset>   8080  TCP
```


Kube-proxy and Cluster IPs

- Cluster Ips are stable virtual IPs,



- `kubectl exec -it $nginx_pod -- env | grep WEBAPP_`

```
PS> kubectl exec -it $nginx_pod -- env | grep WEBAPP_  
WEBAPP_PORT=tcp://10.152.183.44:8080  
WEBAPP_SERVICE_HOST=10.152.183.44  
WEBAPP_PORT_8080_TCP_PROTO=tcp  
WEBAPP_PORT_8080_TCP_ADDR=10.152.183.44  
WEBAPP_SERVICE_PORT=8080  
WEBAPP_PORT_8080_TCP=tcp://10.152.183.44:8080  
WEBAPP_PORT_8080_TCP_PORT=8080
```

- `kubectl delete services,deployments -l run=webapp`

Ingress

Ingress

- Exposing service with http or https
- Other services use:
 - NodePort
 - LoadBalancer
- Uses an ingress controller
 - Nginx-ingress
 - HAProxy, Traefik, Istio, Kong,

Single service

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: test-ingress
spec:
  backend:
    serviceName: testsvc
    servicePort: 80
```

Ingress

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: test-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
    - http:
        paths:
          - path: /testpath
            backend:
              serviceName: test
              servicePort: 80
```

Enable Ingress

```
multipass exec microk8s-vm --  
/snap/bin/microk8s.enable ingress
```

```
Enabling Ingress  
deployment.extensions/default-http-backend unchanged  
service/default-http-backend unchanged  
serviceaccount/nginx-ingress-microk8s-serviceaccount unchanged  
clusterrole.rbac.authorization.k8s.io/nginx-ingress-microk8s-clusterrole unchanged  
role.rbac.authorization.k8s.io/nginx-ingress-microk8s-role unchanged  
clusterrolebinding.rbac.authorization.k8s.io/nginx-ingress-microk8s unchanged  
rolebinding.rbac.authorization.k8s.io/nginx-ingress-microk8s unchanged  
configmap/nginx-load-balancer-microk8s-conf unchanged  
daemonset.apps/nginx-ingress-microk8s-controller unchanged  
Ingress is enabled
```


DaemonSets

- Ensures a copy of a Pod is running across a set of nodes in a Kubernetes cluster
- Typically used for log collectors, monitoring tools etc.
- Like ReplicaSets. DaemonSets run exactly 1 copy per node.
- Add additional capabilities and features to the Kubernetes cluster itself instead of running normal services

DaemonSet definition

```
apiVersion: extensions/v1beta1
kind: DaemonSet
metadata:
  name: fluentd
  namespace: kube-system
  labels:
    app: fluentd
spec:
  template:
    metadata:
      labels:
        app: fluentd
    spec:
      containers:
        - name: fluentd
          image: fluent/fluentd:v0.14.10
```

Node selectors with DaemonSets

- spec:
- nodeSelector:
- ssd: "true"

Jobs

- So far we've looked at long running processes
- Jobs are short-lived, typically one-off tasks
- A Job creates Pods that run until successful termination
- Pods restart even when successful terminated

Job Types

Type	Use case	Behavior	completions	parallelism
One shot	Database migrations	A single pod running once until successful termination	1	1
Parallel fixed completions	Multiple pods processing a set of work in parallel	One or more Pods running one or more times until reaching a fixed completion count	1+	1+
Work queue: parallel Jobs	Multiple pods processing from a centralized work queue	One or more Pods running once until successful termination	1	2+

- `kubectl run -i oneshot \`
- `--image=<some-image:latest> \`
- `--restart=OnFailure`

- `kubectl delete jobs oneshot`

- `kubectl get pod -a -l job-name=oneshot`
- Describe, edit etc are available

```
apiVersion: batch/v1
kind: Job
metadata:
  name: oneshot
  labels:
    chapter: jobs
spec:
  template:
    metadata:
      labels:
        chapter: jobs
    spec:
      containers:
        - name: my-job-container
          image: <some-image>
          imagePullPolicy: Always
          args:
restartPolicy: OnFailure
```

Parallel jobs

- Configure the following variables.
- completions: 10
 - Number of jobs to run
- parallelism: 5
 - Max. parallel executions (same time)


```
apiVersion: batch/v1
kind: Job
metadata:
  name: parallel
  labels:
    chapter: jobs
spec:
  parallelism: 5
  completions: 10
  template:
    metadata:
      labels:
        chapter: jobs
    spec:
      containers:
      - name: kuard
        image: <some-image>
        imagePullPolicy: Always
        args:
restartPolicy: OnFailure
```