


Access control

RBAC

- Role based access control
- Defines who can do what in the cluster
 - Developers
 - Admins
 - Etc.
- Role
 - Verb: get, list..
 - Nouns: Pods, Volumes..



Subjects

ConfigMaps

Pod Service AutoScaler

Deployment Secrets

PV Ingress

ReplicaSets

Namespace DaemonSet

CronJob Job PVC Nodes

API Resources

list get

create watch

delete patch

**Operations
(Verbs)**

RBAC

- Roles need to be connected to users/groups
- 2 types of resources
 - Namespace resources
 - RoleBinding
 - Cluster resources
 - ClusterRoleBinding (entire cluster)

Role example

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: default
  name: pod-reader
rules:
- apiGroups: [""]
  # "" indicates the core API group
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: pod-reader
  namespace: default
subjects:
- kind: User
  name: jane
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role #Role or ClusterRole
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```

ClusterRole

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  # "namespace" omitted since ClusterRoles are not namespaced
  name: secret-reader
rules:
- apiGroups: [""]
  resources: ["secrets"]
  verbs: ["get", "watch", "list"]
```

ClusterRoleBinding

```
apiVersion: rbac.authorization.k8s.io/v1
# This cluster role binding allows anyone in the "manager" group to read secrets
in any namespace.
kind: ClusterRoleBinding
metadata:
  name: read-secrets-global
subjects:
- kind: Group
  name: manager # Name is case sensitive
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: ClusterRole
  name: secret-reader
  apiGroup: rbac.authorization.k8s.io
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: aggregate-cron-tabs-edit
  labels:
    # Add these permissions to the "admin" and "edit" default roles.
    rbac.authorization.k8s.io/aggregate-to-admin: "true"
    rbac.authorization.k8s.io/aggregate-to-edit: "true"
rules:
- apiGroups: ["stable.example.com"]
  resources: ["crontabs"]
  verbs: ["get", "list", "watch", "create", "update", "patch", "delete"]
---
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: aggregate-cron-tabs-view
  labels:
    # Add these permissions to the "view" default role.
    rbac.authorization.k8s.io/aggregate-to-view: "true"
rules:
- apiGroups: ["stable.example.com"]
  resources: ["crontabs"]
  verbs: ["get", "list", "watch"]
```


Service account

```
kubectl create rolebinding my-sa-view \  
  --clusterrole=view \  
  --serviceaccount=my-namespace:my-sa \  
  --namespace=my-namespace
```

if pods don't have a serviceaccount specified, default is used

Users in Kubernetes

- 2 types:
 - service accounts managed by Kubernetes
 - normal users
- Authentication strategies
 - X509 Client Certs
 - Static Token File

X509 Client Certificate

- Add the certificate authority to:
--client-ca-file=SOMEFILE in kube-apiserver configuration
- For MicroK8s:
#kube-apiserver config
sudo vim /var/snap/microk8s/current/args/kube-apiserver
sudo systemctl restart snap.microk8s.daemon -
apiserver.service

Static token file

- Kube-apiserver configuration:
--token-auth-file=SOMEFILE
- File should be CSV, with these columns (min first 3):
token, username, user uid, groups
- Example:
token,user,uid,"group1,group2,group3"
- Requests should have 'Authorization: Bearer <token>' in the headers

Static Password file

- Kube-apiserver config: **--basic-auth-file=SOMEFILE**
- CSV file with:
password,user,uid,"group1,group2,group3"
- Authorization: Basic BASE64ENCODED(USER:PASSWORD)

Service account tokens

- `--service-account-key-file`
A file containing a PEM encoded key for signing bearer tokens. If unspecified, the API server's TLS private key will be used.
- `--service-account-lookup`
If enabled, tokens which are deleted from the API will be revoked.

Service accounts

```
kubectl create serviceaccount nginx  
kubectl get serviceaccounts jenkins -o yaml
```

```
apiVersion: v1
```

```
kind: ServiceAccount
```

```
metadata:
```

```
  # ...
```

```
secrets:
```

```
- name: jenkins-token-1yvwg
```

Service accounts

```
kubectl get secret jenkins-token-1yvwg -o yaml
```

```
apiVersion: v1
```

```
data:
```

```
  ca.crt: (APISERVER'S CA BASE64 ENCODED)
```

```
  namespace: ZGVmYXVsdA==
```

```
  token: (BEARER TOKEN BASE64 ENCODED)
```

```
kind: Secret
```

```
metadata:
```

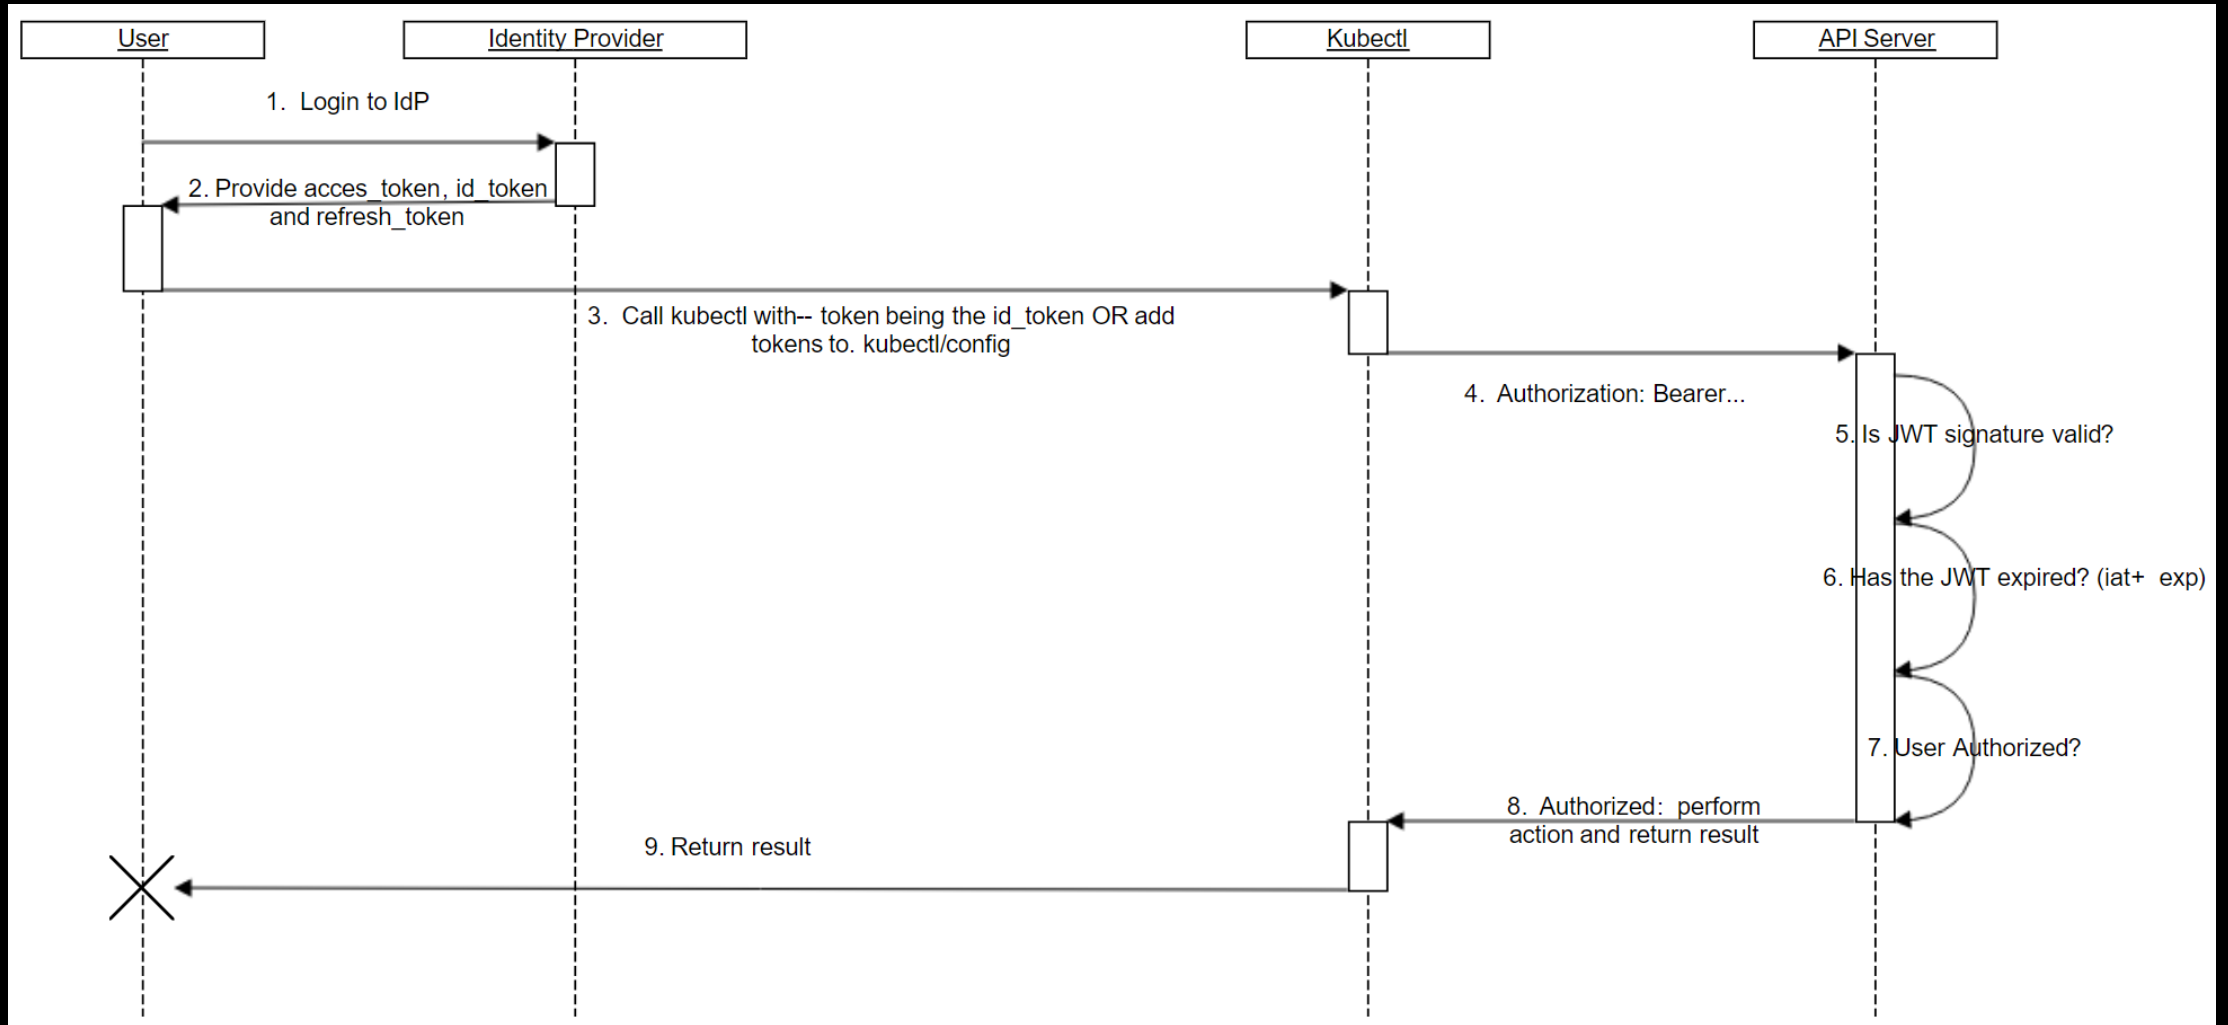
```
  # ...
```

```
type: kubernetes.io/service-account-token
```


OpenID Connect Tokens

- OpenID Connect is a flavor of OAuth2
 - supported by some OAuth2 providers, notably:
 - Azure Active Directory,
 - Salesforce,
 - Google

OpenID Connect Tokens



OpenID Connect

1. Login to your identity provider
2. Your identity provider will provide you with an `access_token`, `id_token` and a `refresh_token`
3. When using `kubectl`, use your `id_token` with the `--token` flag or add it directly to your `kubeconfig`
4. `kubectl` sends your `id_token` in a header called `Authorization` to the API server
5. The API server will make sure the JWT signature is valid by checking against the certificate named in the configuration
6. Check to make sure the `id_token` hasn't expired
7. Make sure the user is authorized
8. Once authorized the API server returns a response to `kubectl`
9. `kubectl` provides feedback to the user

Parameter	Description	Example	Required
<code>--oidc-issuer-url</code>	URL of the provider which allows the API server to discover public signing keys. Only URLs which use the <code>https://</code> scheme are accepted. This is typically the provider's discovery URL without a path, for example " <code>https://accounts.google.com</code> " or " <code>https://login.salesforce.com</code> ". This URL should point to the level below <code>.well-known/openid-configuration</code>	If the discovery URL is <code>https://accounts.google.com/.well-known/openid-configuration</code> , the value should be <code>https://accounts.google.com</code>	Yes
<code>--oidc-client-id</code>	A client id that all tokens must be issued for.	kubernetes	Yes
<code>--oidc-username-claim</code>	JWT claim to use as the user name. By default <code>sub</code> , which is expected to be a unique identifier of the end user. Admins can choose other claims, such as <code>email</code> or <code>name</code> , depending on their provider. However, claims other than <code>email</code> will be prefixed with the issuer URL to prevent naming clashes with other plugins.	sub	No
<code>--oidc-username-prefix</code>	Prefix prepended to username claims to prevent clashes with existing names (such as <code>system: users</code>). For example, the value <code>oidc:</code> will create usernames like <code>oidc:jane.doe</code> . If this flag isn't provided and <code>--oidc-user-claim</code> is a value other than <code>email</code> the prefix defaults to <code>(Issuer URL)#</code> where <code>(Issuer URL)</code> is the value of <code>--oidc-issuer-url</code> . The value <code>-</code> can be used to disable all prefixing.	<code>oidc:</code>	No
<code>--oidc-groups-claim</code>	JWT claim to use as the user's group. If the claim is present it must be an array of strings.	groups	No
<code>--oidc-groups-prefix</code>	Prefix prepended to group claims to prevent clashes with existing names (such as <code>system: groups</code>). For example, the value <code>oidc:</code> will create group names like <code>oidc:engineering</code> and <code>oidc:infra</code> .	<code>oidc:</code>	No
<code>--oidc-required-claim</code>	A key=value pair that describes a required claim in the ID Token. If set, the claim is verified to be present in the ID Token with a matching value. Repeat this flag to specify multiple claims.	<code>claim=value</code>	No
<code>--oidc-ca-file</code>	The path to the certificate for the CA that signed your identity provider's web certificate. Defaults to the host's root CAs.	<code>/etc/kubernetes/ssl/kc-ca.pem</code>	No

Setup kubectl to use OpenID Connect

```
kubectl config set-credentials USER_NAME \  
  --auth-provider=oidc \  
  --auth-provider-arg=idp-issuer-url=( issuer url ) \  
  --auth-provider-arg=client-id=( your client id ) \  
  --auth-provider-arg=client-secret=( your client secret ) \  
  --auth-provider-arg=refresh-token=( your refresh token ) \  
  --auth-provider-arg=idp-certificate-authority=( path to cacert ) \  
  --auth-provider-arg=id-token=( your id_token )
```

Kubernetes Security

```
apiVersion: audit.k8s.io/v1beta1
kind: Policy
omitStages:
  - "RequestReceived"
rules:
  - level: Request
    users: ["admin"]
    resources:
      - group: ""
        resources: ["*"]
  - level: Request
    user: ["system:anonymous"]
    resources:
      - group: ""
        resources: ["*"]
```

