

The background of the slide features a close-up, slightly blurred image of a mechanical dial or gauge. The dial has a circular scale with numbers ranging from 0 to 20, and a central needle pointing towards the 15 mark. The overall image is in a light, desaturated blue-grey tone, giving it a technical and industrial feel.

Access control

RBAC

role based access control

A Role defines an allowed operation (or set of)


Are assigned to accounts

Role

Consists of 2 things, what object-type and the allowed operation

In Kubernetes terms:

- Verb: get, list, create..
- Nouns: Pods, Volumes, Services..



Subjects



API Resources



**Operations
(Verbs)**

RBAC

Roles need to be connected to users/groups

2 types of resources

- Namespace resources
 - RoleBinding
- Cluster resources
 - ClusterRoleBinding (entire cluster)

Role example

```
apiVersion: rbac.authorization.k8s.io/v1
```

```
kind: Role
```

```
metadata:
```

```
  namespace: default
```

```
  name: pod-reader
```

```
rules:
```

```
- apiGroups: [""] # "" indicates the core API group
```

```
  resources: ["pods"]
```

```
  verbs: ["get", "watch", "list"]
```

Role example

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: pod-reader
  namespace: default
subjects:
- kind: User
  name: jane
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role #Role or ClusterRole
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```

ClusterRole

`apiVersion: rbac.authorization.k8s.io/v1`

`kind: ClusterRole`

`metadata:`

`# "namespace" omitted since ClusterRoles are not namespaced`

`name: secret-reader`

`rules:`

`- apiGroups: [""]`

`resources: ["secrets"]`

`verbs: ["get", "watch", "list"]`

ClusterRoleBinding

apiVersion: rbac.authorization.k8s.io/v1

kind: ClusterRoleBinding

metadata:

name: read-secrets-global

subjects:

- kind: Group

name: manager # Name is case sensitive

apiGroup: rbac.authorization.k8s.io

roleRef:

kind: ClusterRole

name: secret-reader

apiGroup: rbac.authorization.k8s.io

Role aggregation

Kubernetes comes with default cluster-roles:

- **cluster-admin** — Allows super-user access to perform any action on any resource
- **admin** — Full admin access, intended to be granted within a namespace
- **edit** — Edit access except for editing RBAC
- **view** — Allows view for most of objects

We don't want to update these, but we can add permissions to them by aggregating.

Role aggregation

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: aggregate-cron-tabs-edit
  labels:
    # Add these permissions to the "admin" and "edit" default roles.
    rbac.authorization.k8s.io/aggregate-to-admin: "true"
    rbac.authorization.k8s.io/aggregate-to-edit: "true"
rules:
- apiGroups: ["stable.example.com"]
  resources: ["crontabs"]
  verbs: ["get", "list", "watch", "create", "update", "patch", "delete"]
```

Role aggregation

```
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: aggregate-cron-tabs-view
  labels:
    # Add these permissions to the "view" default role.
    rbac.authorization.k8s.io/aggregate-to-view: "true"
rules:
- apiGroups: ["stable.example.com"]
  resources: ["crontabs"]
  verbs: ["get", "list", "watch"]
```

Service account

```
kubectl create rolebinding my-sa-view \  
  --clusterrole=view \  
  --serviceaccount=my-namespace:my-sa \  
  --namespace=my-namespace
```

```
# if pods don't have a serviceaccount specified, default is used
```

Users in Kubernetes

2 types:

- service accounts managed by Kubernetes
- normal users

Authentication strategies

- X509 Client Certs
- Static Token File
- Oauth/OpenID provder

X509 Client Certificate

Add the certificate authority to:

--client-ca-file=SOMEFILE in kube-apiserver configuration

For MicroK8s:

```
#kube-apiserver config
```

```
sudo vim /var/snap/microk8s/current/args/kube-apiserver
```

```
sudo systemctl restart snap.microk8s.daemon -  
apiserver.service
```

Static token file

Kube-apiserver configuration:
`--token-auth-file=SOMEFILE`

File should be CSV, with these columns (min first 3):
token, username, user uid, groups

Example:
`token,user,uid,"group1,group2,group3"`

Requests should have 'Authorization: Bearer <token>' in the headers

Static Password file

Kube-apiserver config: **--basic-auth-file=SOMEFILE**

CSV file with:

password,user,uid,"group1,group2,group3"

Authorization: Basic BASE64ENCODED(USER:PASSWORD)

Service account tokens

`--service-account-key-file`

A file containing a PEM encoded key for signing bearer tokens. If unspecified, the API server's TLS private key will be used.

`--service-account-lookup`

If enabled, tokens which are deleted from the API will be revoked.

Service accounts

```
kubectl create serviceaccount nginx  
kubectl get serviceaccounts jenkins -o yaml
```

```
apiVersion: v1
```

```
kind: ServiceAccount
```

```
metadata:
```

```
  # ...
```

```
secrets:
```

```
- name: jenkins-token-1yvwg
```

Service accounts

```
kubectl get secret jenkins-token-1yvwg -o yaml
```

```
apiVersion: v1
```

```
data:
```

```
  ca.crt: (APISERVER'S CA BASE64 ENCODED)
```

```
  namespace: ZGVmYXVsdA==
```

```
  token: (BEARER TOKEN BASE64 ENCODED)
```

```
kind: Secret
```

```
metadata:
```

```
  # ...
```

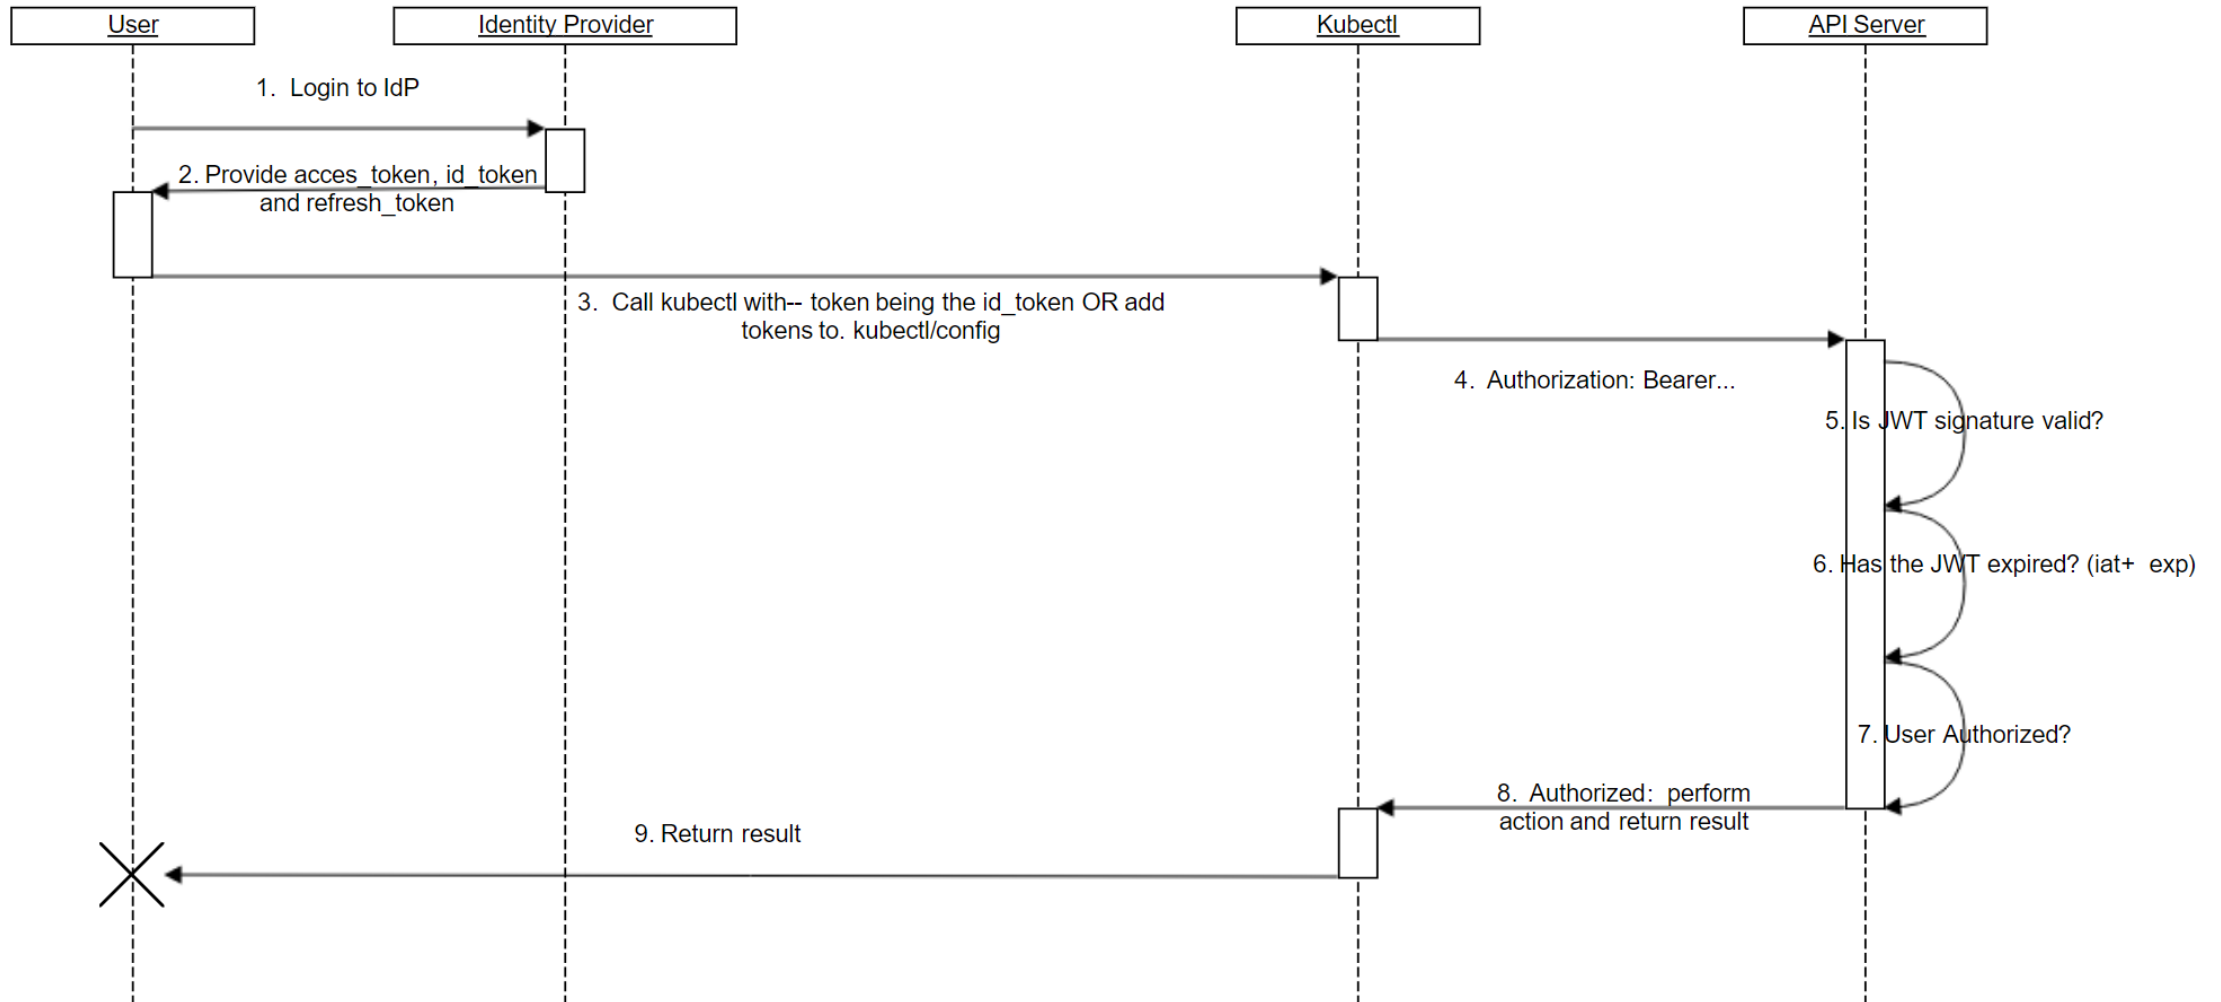
```
type: kubernetes.io/service-account-token
```

OpenID Connect Tokens

OpenID Connect is a flavor of OAuth2

- supported by some OAuth2 providers, notably:
- Azure Active Directory,
- Salesforce,
- Google

OpenID Connect Tokens



OpenID Connect

1. Login to your identity provider
2. Your identity provider will provide you with an `access_token`, `id_token` and a `refresh_token`
3. When using `kubectl`, use your `id_token` with the `--token` flag or add it directly to your `kubeconfig`
4. `kubectl` sends your `id_token` in a header called `Authorization` to the API server
5. The API server will make sure the JWT signature is valid by checking against the certificate named in the configuration
6. Check to make sure the `id_token` hasn't expired
7. Make sure the user is authorized
8. Once authorized the API server returns a response to `kubectl`
9. `kubectl` provides feedback to the user

Parameter	Description	Example	Required
<code>--oidc-issuer-url</code>	URL of the provider which allows the API server to discover public signing keys. Only URLs which use the https:// scheme are accepted. This is typically the provider's discovery URL without a path, for example " https://accounts.google.com " or " https://login.salesforce.com ". This URL should point to the level below .well-known/openid-configuration	If the discovery URL is https://accounts.google.com/.well-known/openid-configuration , the value should be https://accounts.google.com	Yes
<code>--oidc-client-id</code>	A client id that all tokens must be issued for.	kubernetes	Yes
<code>--oidc-username-claim</code>	JWT claim to use as the user name. By default sub , which is expected to be a unique identifier of the end user. Admins can choose other claims, such as email or name , depending on their provider. However, claims other than email will be prefixed with the issuer URL to prevent naming clashes with other plugins.	sub	No
<code>--oidc-username-prefix</code>	Prefix prepended to username claims to prevent clashes with existing names (such as system: users). For example, the value oidc: will create usernames like oidc:jane.doe . If this flag isn't provided and <code>--oidc-user-claim</code> is a value other than email the prefix defaults to (Issuer URL)# where (Issuer URL) is the value of <code>--oidc-issuer-url</code> . The value - can be used to disable all prefixing.	oidc:	No
<code>--oidc-groups-claim</code>	JWT claim to use as the user's group. If the claim is present it must be an array of strings.	groups	No
<code>--oidc-groups-prefix</code>	Prefix prepended to group claims to prevent clashes with existing names (such as system: groups). For example, the value oidc: will create group names like oidc:engineering and oidc:infra .	oidc:	No
<code>--oidc-required-claim</code>	A key=value pair that describes a required claim in the ID Token. If set, the claim is verified to be present in the ID Token with a matching value. Repeat this flag to specify multiple claims.	claim=value	No
<code>--oidc-ca-file</code>	The path to the certificate for the CA that signed your identity provider's web certificate. Defaults to the host's root CAs.	/etc/kubernetes/ssl/kc-ca.pem	No

Setup kubectl to use OpenID Connect

```
kubectl config set-credentials USER_NAME \  
  --auth-provider=oidc \  
  --auth-provider-arg=idp-issuer-url=( issuer url ) \  
  --auth-provider-arg=client-id=( your client id ) \  
  --auth-provider-arg=client-secret=( your client secret ) \  
  --auth-provider-arg=refresh-token=( your refresh token ) \  
  --auth-provider-arg=idp-certificate-authority=( path to cacert ) \  
  --auth-provider-arg=id-token=( your id_token )
```



```
apiVersion: audit.k8s.io/v1beta1
kind: Policy
omitStages:
  - "RequestReceived"
rules:
  - level: Request
    users: ["admin"]
    resources:
      - group: ""
        resources: ["*"]
  - level: Request
    user: ["system:anonymous"]
    resources:
      - group: ""
        resources: ["*"]
```

